

# Deep Learning Basics

## Convolution

HKUST MSBD 6000B

Instructor: Yu Zhang

# Convolutional neural networks

- Strong empirical application performance
- Convolutional networks: neural networks that use convolution in place of general matrix multiplication in at least one of their layers

$$h = \sigma(W^T x + b)$$

for a specific kind of weight matrix  $W$

# Convolution

# Convolution: math formula

- Given functions  $u(t)$  and  $w(t)$ , their convolution is a function  $s(t)$

$$s(t) = \int u(a)w(t - a)da$$

- Written as

$$s = (u * w) \quad \text{or} \quad s(t) = (u * w)(t)$$

# Convolution: discrete version

- Given array  $u_t$  and  $w_t$ , their convolution is a function  $s_t$


$$s_t = \sum_{a=-\infty}^{+\infty} u_a w_{t-a}$$

- Written as


$$s = (u * w) \quad \text{or} \quad s_t = (u * w)_t$$

- When  $u_t$  or  $w_t$  is not defined, assumed to be 0


# Illustration 1



# Illustration 1




# Illustration 1






# Illustration 1: boundary case




# Illustration 1 as matrix multiplication

y	z					a
x	y	z				b
	x	y	z			c
		x	y	z		d
			x	y	z	e
				x	y	f


## Illustration 2: two dimensional case



# Illustration 2



# Illustration 2



# Advantage: sparse interaction

Fully connected layer,  $m \times n$  edges




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Advantage: sparse interaction

Convolutional layer,  $\leq m \times k$  edges




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Advantage: sparse interaction

Multiple convolutional layers: larger receptive field





Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville



# Advantage: parameter sharing



The same kernel are used repeatedly. E.g., the black edge is the same weight in the kernel.




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Advantage: equivariant representations

- Equivariant: transforming the input = transforming the output
- Example: input is an image, transformation is shifting
- $\text{Convolution}(\text{shift}(\text{input})) = \text{shift}(\text{Convolution}(\text{input}))$
- Useful when care only about the **existence** of a pattern, rather than the **location**

Pooling

# Terminology




Figure from *Deep Learning*,  
by Goodfellow, Bengio,  
and Courville

# Pooling

- Summarizing the input (i.e., output the max of the input)




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Motivation from neuroscience

- David Hubel and Torsten Wiesel studied early visual system in human brain (V1 or primary visual cortex), and won Nobel prize for this
- V1 properties
  - 2D spatial arrangement
  - Simple cells: inspire convolution layers
  - Complex cells: inspire pooling layers

Variants of convolution and pooling

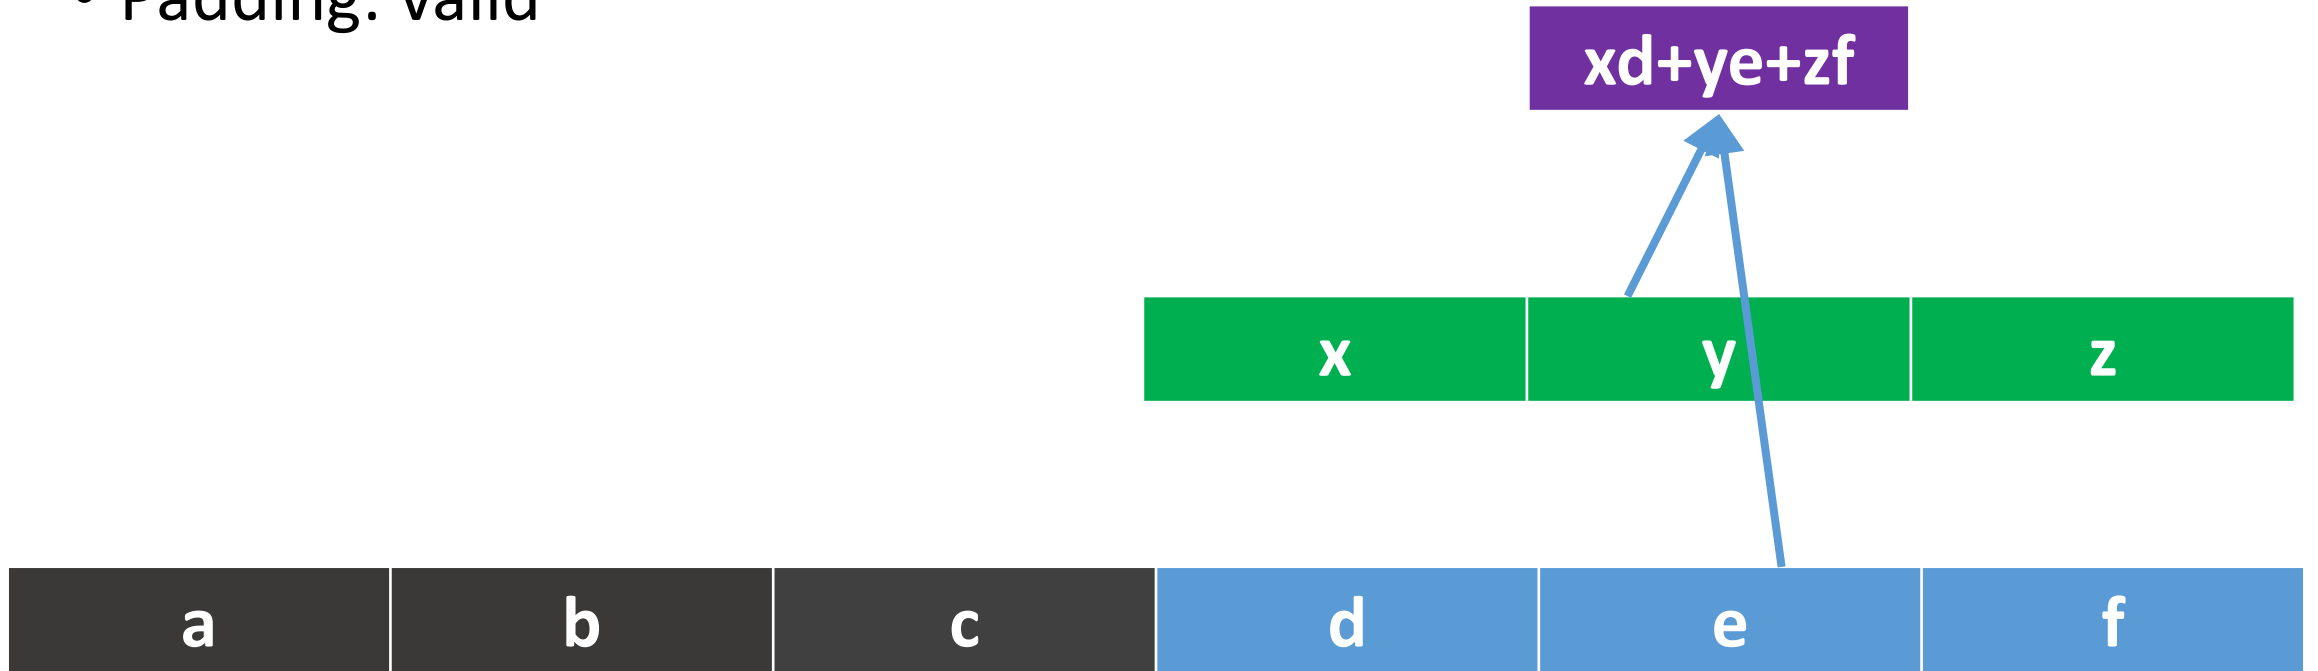
# Variants of convolutional layers

- Multiple dimensional convolution
- Input and kernel can be 3D
  - E.g., images have (width, height, RGB channels)
- Multiple kernels lead to multiple feature maps (also called channels)
- Mini-batch of images have 4D: (image\_id, width, height, RGB channels)




# Variants of convolutional layers

- Padding: valid



# Variants of convolutional layers

- Padding: same



# Variants of convolutional layers

- Stride




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Variants of pooling

- Stride and padding




Figure from *Deep Learning*, by Goodfellow, Bengio, and Courville

# Variants of pooling

- Max pooling  $y = \max\{x_1, x_2, \dots, x_k\}$
- Average pooling  $y = \text{mean}\{x_1, x_2, \dots, x_k\}$
- Others like max-out